


Overview: GCN training and testing (inductive)

Reminder

Learn feature embeddings (DGL-2)

A generalized message propagation rule to update the embeddings in a GCN layer is based on two steps:

- (1) aggregate information from "myself" (a given node) and "my neighbors" (its neighbors), and
- (2) apply a linear transformation. $(\rightarrow \mathcal{L})$ = learnable transformation / supports




The node's neighborhood defines a computation graph over which we aggregate and transform embeddings.


Our ultimate goal: Use the learned embeddings in the final layer for downstream learning tasks (e.g., regression/classification).

Graph learning tasks (DGL-1)

node classification



edge/link prediction




GCN training and its powerful inductive capability at testing

Note: we will examine supervised node- and graph-based classification/regression. In a similar way, we can adopt such formalization and steps to other learning tasks.

Design your GCN layer and optimize your loss function



- 1) Define a neighborhood aggregation function. $\rightarrow$ "active research area"
- 2) Define the building blocks (operations) of a GCN layer (e.g., linear transformation, nonlinear activation function).
- 3) Define your loss function to optimize.
- 4) Define the prediction head (node, edge, or graph), which may include additional hyperparameters (e.g., β and ω encoding the classifier bias and weights to learn).
- 5) Optimize the loss over embeddings and additional hyperparameters (e.g., classifier weights and (biases) during training on the training samples using stochastic gradient descent (SGD).

- Compute the loss during the forward pass.
- Compute the gradient of the loss using the chain rule during the backward pass.
- Update the network (model) parameters. $\theta_t = \theta_{t-1} - \eta \nabla \mathcal{L}(\theta_{t-1})$




Reminder (DGL-2): Through this optimization, we aim to find an ultimate low-dimensional representation space (a much simpler manifold) where linear models can work well on the final embeddings.


Test on unseen nodes



(Reminder) Non-linear multi-regression models



Forward pass through the GNN layers and prediction head





Overview: GCN training and testing (inductive)


GCN training and its powerful inductive capability at testing

Inductive capability of GCNs

What happens if we don't connect the node? Can we still predict its output scores?



Can GCN generalize to many new nodes as they arrive?



One of the most compelling aspects of GNN models lies in their **inductive capability** where we can easily generalize to new graphs with varying sizes / properties.

This power stems from the **feature aggregation** and **parameter sharing** at the node level.


What happens if we don't share the parameters across all nodes?

We could learn a model with **separate parameters** associated with each node. However, the network must **relearn** the meaning of the connections in the graph at each position, and training would require many graphs with the **same topology**.



Can we classify graphs with different sizes?




Using ML or DL can we multi-regress this dataset of multiresolution brain graphs?



How can we extend the node-level regression to a graph-level regression?

Hint: leverage the power of **integration**.

integrate node embeddings across the whole graph and use the integrated embedding to solve the same optimization problem for node-based learning.



Graph-level prediction head: How to combine the node embeddings across the entire graph?

Global mean pooling

$$h_K^G = \text{Mean}\{h_K^{(i)} \in \mathbb{R}^{d_K}, \forall i \in V_G\}$$

average activation at each embedding dimension across nodes.

Global max pooling

$$h_K^G = \text{Max}\{h_K^{(i)} \in \mathbb{R}^{d_K}, \forall i \in V_G\}$$

Keep the strong activation.

Global sum pooling

$$h_K^G = \text{Sum}\{h_K^{(i)} \in \mathbb{R}^{d_K}, \forall i \in V_G\}$$

enhance without normalization.

Global pooling over a (large) graph will lose information and cannot differentiate between different graph structures.

Example: Let us find a case where two graphs G_1 and G_2 have different embedding distributions but their sum produces the same number (e.g., zero).

$$G_1 = [-0.5, 1, -0.6, 0.1] \text{ and } G_2 = [18, 0, -5, -13]$$

We cannot differentiate between G_1 and G_2 using global sum pooling.

Hierarchical pooling (DiffPool 2018)

This allows for a better graph **differentiation**. Think about a toy example that shows this.

Hierarchical Graph Representation Learning with Differentiable Pooling




Figure 1: High-level illustration of our proposed method **DiffPool**. At each **hierarchical layer**, we run a GNN model to obtain embeddings of nodes. We then use these learned embeddings to cluster nodes together and run another GNN layer on this coarsened graph. This whole process is repeated for L layers and we use the final output representation to classify the graph.

Ying, Zhiao, et al. "Hierarchical graph representation learning with differentiable pooling." *Advances in neural information processing systems* 31 (2018). https://proceedings.neurips.cc/paper_files/paper/2018/file/c77dbaf6759f53376604656903637-Paper.pdf


GNN layer operations (building blocks)

1) The operation order can change depending on the rationale and target task and a few operations can be merged or removed.

! 2) How and when to perform an operation? And why?
Example: how and when to "aggregate" and "transform"? We can first transform then aggregate or the other way around.




GNN layer operations (building blocks)



More local and global aggregation methods (UDL-13)

How to aggregate information locally (at the node level) or globally (at the graph level)?



Mean aggregation

$$\mathbf{h}_{k+1}^{(n)} = \mathbf{a} [\beta_k \mathbf{1} + \Omega_k \cdot \mathbf{h}_k^{(n)} + \Omega_k \cdot \text{agg}[n, k]]$$

summing the embeddings of neighbors

$$\text{agg}[n] = \sum_{m \in \text{ne}[n]} \mathbf{h}_m$$

averaging the embeddings of neighbors

$$\text{agg}[n] = \frac{1}{|\text{ne}[n]|} \sum_{m \in \text{ne}[n]} \mathbf{h}_m$$




The new GCN layer can be written using the degree matrix:

$$\mathbf{H}_{k+1} = \mathbf{a} [\beta_k \mathbf{1}^T + \Omega_k \mathbf{H}_k (\mathbf{A} \mathbf{D}^{-1} + \mathbf{I})]$$

Kipf normalization

The logic behind this: information coming from nodes with a very large number of neighbors should be downweighted since there are many connections and they provide less unique information:



$$\text{agg}[n] = \sum_{m \in \text{ne}[n]} \frac{\mathbf{h}_m}{\sqrt{|\text{ne}[n]| |\text{ne}[m]|}}$$

The new GCN layer can be written as:

$$\mathbf{H}_{k+1} = \mathbf{a} [\beta_k \mathbf{1}^T + \Omega_k \mathbf{H}_k (\mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} + \mathbf{I})]$$

Max pooling

Max-pooling can also be applied at the node level to update the node embedding. The operator $\max(\cdot)$ returns the element-wise maximum of the embeddings that are neighbors to node n .

$$\text{agg}[n] = \max_{m \in \text{ne}[n]} [\mathbf{h}_m]$$

Aggregation by attention

Reminder: Previous aggregation methods consider the neighbors equally or in a way that depends on the graph topology $\rightarrow$ the aggregation depends on the node embeddings and local topology.

Rule: first transform, then pay attention.

1) Linearly transform the embeddings of a current node as follows:

$$\mathbf{H}'_k = \beta_k \mathbf{1}^T + \Omega_k \mathbf{H}_k$$




Figure 13.12 Comparison of graph convolutional network, dot product attention, and graph attention network. In each case, the mechanism maps N embeddings of size D stored in a $D \times N$ matrix $\mathbf{X}$ to an output of the same size. The graph convolutional network applies a linear transformation $\mathbf{X}' = \Omega \mathbf{X}$ to the data matrix. It then computes a weighted sum of the transformed data, where the weighting is based on the adjacency matrix. A bias β is added and the result is passed through an activation function. b) The outputs of the self-attention

2) Define a node-to-node similarity matrix that will act as an attention matrix in the embedding aggregation step:


$$s_{mn} = \left[\phi_k^T \begin{bmatrix} \mathbf{h}'_m \\ \mathbf{h}'_n \end{bmatrix} \right] \rightarrow \mathbf{S}$$

As in dot-product self-attention, the attention weights that contribute to each output embedding are normalized to be positive and sum to one using the softmax operation. These weights are applied to the transformed embeddings.

$$\mathbf{H}_{k+1} = \mathbf{a} [\mathbf{H}'_k \cdot \text{Softmax}[\mathbf{S}, \mathbf{A} + \mathbf{I}]]$$

The function $\text{Softmax}(\cdot, \cdot)$ applies the softmax operation separately to each column of its first argument, setting values where the second argument is zero to negative infinity. This has the effect of ensuring that the attention to non-neighboring nodes is zero.

The attentions are masked so that each node only attends to itself and its neighbors.



The graph attention network combines both of these mechanisms; the weights are both computed from the data (i.e., embeddings) and based on the adjacency matrix (graph topology).